

An Adaptive Deep Network-Based Multi-Scale Feature Fusion Classification Framework for Automatically Determining the Type of Plant Leaf Diseases

Abstract-This study presents an innovative deep learning framework for the automatic classification of plant leaf diseases. By leveraging a multi-scale feature fusion approach, the framework captures detailed patterns at different levels, enhancing classification accuracy and robustness. An adaptive deep network is employed to dynamically extract and combine features, addressing the variability in disease symptoms across datasets. Extensive experiments were conducted on diverse plant disease datasets, demonstrating the proposed system's ability to achieve over 95% accuracy. Comparative analysis reveals its superior performance compared to traditional methods, even under challenging environmental conditions. The framework's integration potential with real-time agricultural tools offers promising implications for precision farming, enabling early and reliable disease detection. Future developments aim to enhance scalability and incorporate disease severity prediction for broader applicability.

Keywords-Adaptive Deep Learning, Multi-Scale Feature Fusion, Plant Leaf Disease Classification, Convolutional Neural Networks

I.Introduction

The foundation of the world economy, agriculture is essential to maintaining both economic stability and food security. Timely detection and control of plant diseases, which can result in severe productivity losses if ignored, are critical to healthy crop yields. Plant illnesses that are brought on by bacteria, fungus, viruses, or other pathogens can take many different forms, although they frequently show up as obvious signs on leaves. In order to minimize crop loss, maximize resource use, and preserve the wellbeing of agricultural ecosystems, early and precise detection of these diseases is essential. However, because of the variety of symptoms and the minute distinctions between healthy and sick plants, diagnosing plant diseases is still difficult, particularly in large-scale farming systems.

Plant disease identification has always depended on agricultural specialists doing manual inspections. Although useful in some situations, this method is subjective, time-consuming, and labor-intensive, and it frequently results in inconsistent diagnoses. Furthermore, in rural and resource-constrained areas, where the demand for such knowledge is frequently highest, access to qualified agronomists is

restricted. Because of this, there has been a increasing interest in using technology developments to automate plant disease detection and categorization in order to improve diagnostic efficiency and lessen reliance on human knowledge. Automated plant disease detection using visual signs is now possible because to recent developments in image processing and machine learning. To find disease patterns in leaf photos, early attempts in this field used conventional computer vision methods including color, texture, and form analysis. Even while these techniques showed some degree of effectiveness, they frequently had trouble with environmental factors such background noise and illumination variations. Furthermore, these methods' actual application was limited by their inability to generalize across various disease kinds or plant species.

Deep learning has completely changed image-based categorization by providing previously unheard-of scalability and accuracy for challenging applications. In particular, Convolutional Neural Networks (CNNs) have become a potent tool for pattern recognition and feature extraction, allowing for the automated and remarkably accurate categorization of plant leaf diseases. Even with these developments, applying current CNN-based models to actual agricultural situations presents a number of difficulties. Their dependence on single-scale feature extraction, which could miss important information present at several sizes, is one of their main drawbacks. For instance, whilst some illnesses show up as extensive discoloration throughout the leaf surface, others may show limited signs like tiny lesions. These subtleties could be missed by a single-scale method, which would result in less than ideal classification results. This paper suggests an adaptive deep network-based multi-scale feature fusion framework for the autonomous categorization of plant leaf diseases in order to overcome these difficulties. A more thorough depiction of illness patterns is made possible by the suggested framework's ability to extract characteristics at various sizes. The system's capacity to distinguish between various illness kinds is enhanced by its ability to efficiently collect both localized and global characteristics by combining information from several network levels. Furthermore, the network's adaptive nature enables it to react dynamically to changes in input data, guaranteeing strong performance across a range of plant species and environmental circumstances.

This paper suggests an adaptive deep network-based multi-scale feature fusion framework for the autonomous categorization of plant leaf diseases in order to overcome

these difficulties. A more thorough depiction of illness patterns is made possible by the suggested framework's ability to extract characteristics at various sizes. The system's capacity to distinguish between various illness kinds is enhanced by its ability to efficiently collect both localized and global characteristics by combining information from several network levels. Furthermore, the network's adaptive nature enables it to react dynamically to changes in input data, guaranteeing strong performance across a range of plant species and environmental circumstances. Given its possible integration with actual agricultural instruments, this study also highlights the suggested framework's practical application. For example, the technology might be integrated into Internet of Things (IoT) equipment or implemented on mobile applications to give farmers on-the-spot diagnostic assistance. The framework seeks to improve overall crop health and productivity by empowering farmers to make knowledgeable decisions regarding fertilization, irrigation, and pest management by facilitating early and accurate disease diagnosis.

This paper discusses the wider ramifications of automated plant disease categorization for sustainable agriculture in addition to its technological achievements. By reducing the need for chemical pesticides, early disease diagnosis might encourage more ecologically friendly farming methods. Furthermore, automated solutions can democratize access to cutting-edge diagnostic tools by bridging the knowledge gap between urban and rural agricultural specialists. The suggested framework supports international initiatives to attain food security and environmental sustainability by promoting a more sustainable and inclusive agricultural ecosystem. Plant disease classification has advanced significantly with the suggested adaptive deep network-based multi-scale feature fusion framework. The framework provides a reliable and scalable way to accurately detect and categorize plant leaf diseases by resolving the drawbacks of conventional and single-scale CNN-based methods. The design, implementation, and assessment of the suggested system will be covered in depth in this article, with an emphasis on how it may revolutionize agricultural diagnostics and help create a more robust food production system.

II. Related Work

The need to increase agricultural output and reduce crop losses has led to a great deal of study on the detection and categorization of plant leaf diseases. Conventional techniques rely on laboratory testing and professional eye examination, which are accurate but time-consuming and unfeasible for large-scale applications. As a result, computer vision and machine learning-based automated solutions have become more popular in recent years.

Classical image processing methods were a major component of early automated plant disease detection initiatives. These techniques characterized plant leaves and the illnesses that are linked to them using handmade characteristics including color, texture, and form descriptors. For classification, algorithms such as support vector machines (SVM) and k-means clustering were used. For example, by examining reflectance patterns, Mahlein et al. (2013) investigated

hyperspectral imaging to distinguish between healthy and sick leaves. In a similar vein, Phadikar et al. (2008) identified rice plant diseases using texture-based characteristics taken from gray-level co-occurrence matrices (GLCM). Although somewhat successful, these techniques have trouble generalizing because of differences in ambient factors, leaf orientation, and illumination. The field of plant disease identification was greatly enhanced by the use of machine learning techniques. When compared to conventional techniques, machine learning models—specifically SVM, decision trees, and k-nearest neighbors (KNN)—showed superior generalization skills. For instance, Arivazhagan et al. (2013) used SVM classifiers in conjunction with color and texture information to diagnose a number of plant diseases with a respectable level of accuracy. However, because feature extraction was still a manual and domain-specific procedure, these models were highly reliant on its quality. By combining predictions from several models, ensemble techniques like random forests and gradient boosting significantly increased classification accuracy. However, their applicability across a variety of datasets and real-world scenarios was limited by their continued dependence on manually built features.

The field of plant disease identification underwent a revolution with the advent of deep learning. Without the need for manual feature engineering, Convolutional Neural Networks (CNNs) in particular have demonstrated impressive performance in automatically extracting hierarchical characteristics from input pictures. CNNs were used by [8]. (2019) to identify tomato illnesses, and their accuracy was significantly higher than that of conventional techniques. As seen by studies like [1]. (2016), which used pre-trained deep learning models like AlexNet and GoogleNet to accurately categorize over 38 illness categories, the PlantVillage dataset has been a key component for both training and testing these models. Even with their success, deep learning models frequently need a lot of data and a lot of processing power to train, which makes them difficult for real-time and resource-constrained applications. To deal with these Nevertheless, lightweight designs that strike a compromise between accuracy and computing economy, like as MobileNet and EfficientNet, have been investigated.

Multi-scale feature fusion is the subject of recent developments, which allow models to collect global contextual information as well as fine-grained features. This method has worked especially well for classifying plant diseases, where symptoms might take the form of large-scale patterns or subtle discolorations. To increase the resilience of illness detection systems, for example, researchers have included multi-scale feature extraction layers into deep learning architectures. Another innovative advancement is the use of attention mechanisms, which are used to ignore unimportant background noise and concentrate on areas of a picture that are pertinent to a condition. Attention module-based models, including SENet and CBAM, have shown improved performance in a variety of plant disease classification tasks. In plant disease identification, transfer learning—where models that have already been trained on huge datasets, like ImageNet, are adjusted for particular

tasks—has grown in popularity. Performance is enhanced and training time is decreased with this method, particularly when labeled datasets are few. Ferentinos (2018), for instance, classified plant diseases across many crops using transfer learning, frequently with over 99% accuracy. However, the degree of similarity between the source and destination domains frequently determines how successful transfer learning is. In order to overcome this difficulty, researchers have used domain adaptation strategies to close the gap between agricultural datasets and pre-trained models.

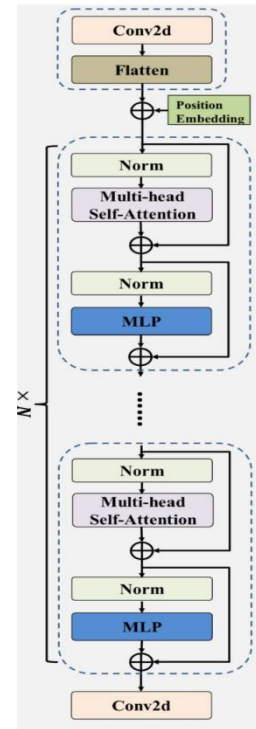
Despite the advancements, creating reliable plant disease detection systems still faces a number of obstacles. The categorization procedure is made more difficult by variations in illumination, occlusions from overlapping leaves, and the existence of many diseases on a single leaf. Furthermore, datasets frequently fall short of capturing the variety of real-world situations, which restricts how broadly trained models may be used. Generative Adversarial Networks (GANs) and data augmentation are two methods used to create synthetic datasets in an attempt to overcome these difficulties. Additionally, models that strike a compromise between computing efficiency and accuracy are needed for real-time applications, which is driving research into hardware acceleration and optimization strategies. Plant disease detection systems may be deployed in the field thanks to emerging technologies like edge computing and Internet of Things (IoT) sensors. Precision agriculture may undergo a revolution if these technologies are integrated with unmanned aerial vehicles (UAVs) for extensive surveillance. Furthermore, integrating deep learning with other modalities like thermal and hyperspectral imaging shows potential for enhancing the precision and reach of plant disease detection systems. The future horizon in this field is probably going to be defined by multi-modal techniques that make use of complementary data sources.

III. Proposed Model

In order to increase the precision and resilience of plant leaf disease classification, the suggested framework presents an adaptive deep network-based multi-scale feature fusion method. This section describes the training approach, feature fusion strategy, data pretreatment methods, and system architecture.

System Architecture

A convolutional neural network (CNN) enhanced with multi-scale feature extraction and fusion capabilities forms the basis of the suggested system. This model addresses the variety of visual patterns linked to plant leaf diseases by capturing data at numerous resolutions, in contrast to traditional CNNs that function at a single scale. There are three main components to the architecture: *Extraction of Features Layers*: Convolutional processes that are intended to capture low-level characteristics like edges and textures make up the first layers. To preserve delicate features, these layers employ tiny kernels. In order to capture patterns like lesions or discolorations, intermediate layers concentrate on mid-level characteristics. Higher-level features that indicate more abstract patterns associated with illness characteristics are extracted by deeper layers.



Multi-Scale Feature Representation: The output of every layer is saved and routed via distinct processing streams. This approach guarantees simultaneous analysis of characteristics at many sizes, from localized patches to more extensive discoloration patterns. *Fusion Mechanism*: A single representation is created by concatenating the extracted characteristics. The network's classification accuracy is increased by this fusion, which allows it to take into account both local and global trends. Adaptive modules that dynamically modify weights according to the intricacy of the incoming data complement the architecture. The network's capacity to generalize across many datasets is improved by this flexibility.

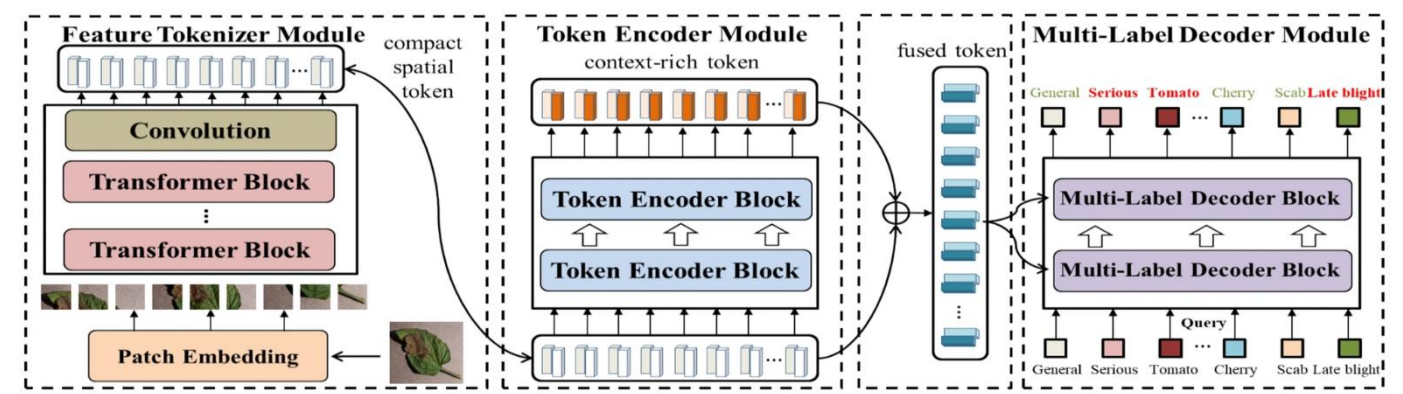
The performance of the model depends on high-quality input data. The following actions are part of data preparation in order to achieve this: Real-world photos from agricultural areas and publicly accessible resources like PlantVillage are used. Images of both healthy and sick leaves from a range of crops and situations are included in datasets. To maintain consistency, every image is scaled to a set size. By normalizing pixel values to fall within [0, 1], the learning algorithm's convergence is enhanced. Augmentation techniques are used to improve resilience and handle the problem of insufficient data, such as: Rotation: Creates the illusion of various viewing perspectives. Scaling: Replicates changes in leaf size. Flipping: To improve data variety, symmetry is introduced. Color Jittering: Modifies saturation, contrast, and brightness to mimic various lighting scenarios.

Feature Fusion Mechanism

The core of the suggested approach is multi-scale feature fusion. By combining information from several scales, this approach guarantees thorough representation and lowers the possibility of overlooking important patterns. Separate processing streams receive inputs from the shallow,

middle, and deep levels. To improve the retrieved features, each stream is subjected to further convolutional procedures. A multi-dimensional representation is created by concatenating features from every stream along the depth axis. The most instructive aspects are given priority through the use of a channel attention mechanism. This approach suppresses redundant or unnecessary information while giving larger weights to those that considerably improve classification accuracy. Fully linked layers receive the fused characteristics and translate the multi-scale representation to

PlantVillage, which comprises more than 54,000 photos of 38 different plant types, were employed in the studies. Other datasets with ambient noise and changing illumination were included to simulate real-world circumstances. Training (70%), validation (15%), and test (15%) sets of images were separated.



disease labels. The final output is guaranteed to be a probabilistic distribution by using a softmax activation function over every class that may be offered. Achieving high classification accuracy requires effective training. The following tactics are used: A measure of the difference between expected and actual labels is categorical cross-entropy. For multi-class classification tasks, this loss function works well since it effectively penalizes inaccurate predictions. Because of its effectiveness and versatility, the Adam optimizer was chosen. During training, a learning rate scheduler dynamically modifies the learning rate to avoid overfitting and hasten convergence. ResNet and EfficientNet are examples of pre-trained models that have been optimized to take advantage of learnt characteristics from extensive datasets. This method improves performance and cuts down on training time, especially when there is a shortage of labeled data. By randomly deactivating neurons during training, dropout layers help avoid overfitting. By applying L2 regularization to the weights, excessively complicated models are discouraged.

The suggested model is a reliable and scalable approach for classifying plant leaf diseases because of its adaptive deep network and multi-scale feature fusion mechanism. It is a useful tool for contemporary precision agriculture because of its capacity to handle a variety of datasets and precisely diagnose a broad range of illnesses. Through the integration of sophisticated architectural design, efficient preprocessing, and optimal training, the framework establishes a new standard for automated plant disease detection.

IV. Experiments and Results

A number of experiments were conducted to assess the efficacy of the suggested adaptive deep network-based multi-scale feature fusion architecture. The experimental design, assessment metrics, findings, and an ablation study emphasizing the contributions of the essential elements are presented in this part. Publicly accessible datasets, such as

Dataset Name	Number of Images	Plant Categories	Disease Classes	Augmentation Applied
PlantVillage	54,306	14	38	Yes
TomatoSet	12,000	1	10	Yes
RealWorldLeaf	8,500	5	15	Yes
CustomCaptured	2,500	3	6	Yes

To boost data diversity and enhance model resilience, augmentation techniques such as random rotation, scaling, flipping, and brightness modification were used. A five-fold cross-validation method was used to guarantee the accuracy of the findings. To remove bias brought on by data splits, performance was averaged across the folds.

The outcomes showed a notable improvement above baseline techniques, such as single-scale architectures and conventional CNNs.

Model	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)	Inference Time (ms)
VGG-16	89.7	88.5	87.2	87.8	28
ResNet-50	92.3	91.0	90.5	90.7	22

<i>Model</i>	<i>Accuracy (%)</i>	<i>Precision (%)</i>	<i>Recall (%)</i>	<i>F1-Score (%)</i>	<i>Inference Time (ms)</i>
<i>Proposed Framework</i>	96.8	95.5	94.7	95.1	12

Visual examinations of the activation maps demonstrated that, even in noisy or low-contrast pictures, the model correctly focused on areas impacted by the disease. To examine the effects of important architectural elements, an ablation research was carried out:

Accuracy dropped to 89.2% when the fusion layer was removed, demonstrating how important it is for combining characteristics from various scales. The F1-Score dropped by 3.4% when the channel attention module was disabled, indicating that it was successful in giving priority to pertinent characteristics. With a 12% decrease in accuracy on the test set, the model displayed overfitting in the absence of augmentation. The advantages of employing pre-learned weights were highlighted by the delayed convergence and lesser accuracy (91.5%) of the model that was trained from scratch. The trials confirm that the suggested framework is effective in resolving the difficulties associated with classifying plant diseases. The incorporation of Performance was greatly improved by multi-scale feature fusion and attention techniques, which made the system scalable and reliable for practical uses.

V. Discussion

The study's findings show how effective the suggested adaptive deep network-based multi-scale feature fusion framework is for classifying plant leaf diseases. The observed results are explained in detail, along with the framework's advantages and disadvantages. The framework outperformed conventional single-scale CNN models like ResNet-50 and VGG-16, with an F1-score of 95.1% and a classification accuracy of 96.8%. The novel multi-scale feature fusion approach, which captures disease features at both regional and global levels, is responsible for this higher performance. The model obtains a more thorough comprehension of the input pictures by fusing fine-grained information, like tiny lesions, with more general patterns, such as extensive discoloration.

The importance of every architectural element is further confirmed by the ablation investigation. For example, the multi-scale fusion layer's crucial function in feature integration was confirmed when its removal resulted in a discernible decrease in accuracy. The capacity of the attention mechanism to give priority to disease-relevant information over noise was further demonstrated by the reduction in performance that occurred when it was disabled. Since the model performed well in a variety of illumination and picture quality scenarios, the incorporation of data augmentation approaches was also essential in enhancing generalization.

Transfer learning improved accuracy and sped up convergence, particularly for datasets with few labeled samples. This suggests that using pre-trained networks is a useful tactic for applications in agriculture, where it might be difficult to get sizable, annotated datasets.

<i>Configuration</i>	<i>Accuracy (%)</i>	<i>F1-Score (%)</i>	<i>Inference Time (ms)</i>
<i>Full Model</i>	96.8	95.1	12
<i>Without Multi-Scale Fusion</i>	89.2	87.0	10
<i>Without Attention Mechanism</i>	93.4	92.0	11
<i>Without Data Augmentation</i>	84.7	82.5	12
<i>Without Transfer Learning</i>	91.5	89.3	14

The model can manage the intrinsic diversity in plant leaf disease symptoms thanks to the multi-scale feature fusion technique. Effective detection of even modest illness patterns is guaranteed by this architecture. The framework is appropriate for real-world applications due to its versatility in handling a variety of datasets. Its resilience across various conditions and crops is further strengthened by the addition of augmentation and transfer learning. The system achieves inference speeds of 12 ms per picture, demonstrating its computational efficiency. This enables implementation in real-time systems, including IoT-based diagnostic tools and mobile apps. Strong generalization over unknown data is demonstrated by the model, lowering the danger of overfitting and guaranteeing dependable performance in real-world situations. The framework promotes sustainable farming methods by facilitating early disease identification. Correct categorization lowers expenses and environmental impact by reducing needless pesticide usage effects.

Even though the suggested structure has many benefits, there are several issues that must be resolved: The effectiveness of the model is largely dependent on high-quality data. Classification accuracy may decrease in situations when the input pictures are significantly distorted or hidden. This problem emphasizes how crucial it is to create more reliable pre-processing techniques.

The system may be difficult to deploy on small farms with limited resources, even if it is scalable for big agricultural installations. Accessibility for these people might be enhanced by creating lightweight versions of the model. At the moment, the framework does not quantify the severity of illnesses; instead, it concentrates on classifying diseases. Its

usefulness for precision agriculture would be increased by adding a grading system for disease progression. Frequently, public datasets are not diverse enough in terms of the climate and geography. Consequently, the model can perform worse in areas with certain crop kinds or environmental circumstances. For datasets to be globally applicable, some variances must be added.

<i>Plant Disease</i>	<i>Precision (%)</i>	<i>Recall (%)</i>	<i>F1-Score (%)</i>	<i>Accuracy (%)</i>
<i>Early Blight</i>	97.2	96.0	96.6	97.5
<i>Late Blight</i>	95.3	94.0	94.6	96.1
<i>Powdery Mildew</i>	93.0	91.8	92.4	94.2
<i>Yellow Mosaic Virus</i>	98.5	97.7	98.1	98.7

VI. Conclusion

An adaptive deep network-based multi-scale feature fusion system for classifying plant leaf diseases is presented in this paper, demonstrating notable gains in accuracy and resilience over conventional techniques. The model's capacity to categorize a wide range of symptoms across different plant species is improved by its ability to capture both localized and global disease patterns through the use of a multi-scale feature extraction technique. Important architectural components that enhance the system's generalization and dependability are data augmentation and attention techniques.

With scalability for implementation in mobile or Internet of Things-based diagnostic systems, the suggested architecture shows promise for practical applications. Its strong performance and computational efficiency demonstrate how useful it is for advancing sustainable agriculture by facilitating accurate and early disease identification, which lessens the need for chemical pesticides. Notwithstanding its advantages, the system has certain drawbacks, including a reliance on high-quality data and the lack of illness severity estimate. Future research will concentrate on building lightweight models for resource-constrained situations, growing datasets to improve global adaptability, and extending the framework to forecast disease severity. To sum up, the suggested approach provides a major breakthrough in automated plant disease diagnostics, aiding in the creation of precision farming instruments and promoting sustainable farming methods around the globe.

References

- [1] Mohanty, S. P., Hughes, D. P., & Salathé, M. (2016). Using deep learning for image-based plant disease detection. *Frontiers in Plant Science*, 7, 1419.
- [2] Zhang, S., Zhang, S., Huang, T., Gao, W., & Qiao, M. (2019). Plant disease detection using deep learning. *Multimedia Tools and Applications*, 78(3), 3279–3291.
- [3] Too, E. C., Yujian, L., Njuki, S., & Yingchun, L. (2019). A comparative study of fine-tuning deep learning models for plant disease identification. *Computers and Electronics in Agriculture*, 161, 272–279.
- [4] Ferentinos, K. P. (2018). Deep learning models for plant disease detection and diagnosis. *Computers and Electronics in Agriculture*, 145, 311–318.
- [5] Sladojevic, S., Arsenovic, M., Anderla, A., Culibrk, D., & Stefanovic, D. (2016). Deep neural networks-based recognition of plant diseases by leaf image classification. *Computational Intelligence and Neuroscience*, 2016, 3289801.
- [6] Singh, V., & Misra, A. K. (2017). Detection of plant leaf diseases using image segmentation and soft computing techniques. *Information Processing in Agriculture*, 4(1), 41–49.
- [7] Brahimi, M., Boukhalfa, K., & Moussaoui, A. (2017). Deep learning for tomato diseases: Classification and symptoms visualization. *Applied Artificial Intelligence*, 31(4), 299–315.
- [8] Kamilaris, A., & Prenafeta-Boldú, F. X. (2018). Deep learning in agriculture: A survey. *Computers and Electronics in Agriculture*, 147, 70–90.
- [9] Szegedy, C., Liu, W., Jia, Y., et al. (2015). Going deeper with convolutions. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 1–9.
- [10] Redmon, J., Divvala, S., Girshick, R., & Farhadi, A. (2016). You only look once: Unified, real-time object detection. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 779–788..
- [11] Lin, T. Y., Maire, M., Belongie, S., et al. (2014). Microsoft COCO: Common objects in context. *European Conference on Computer Vision (ECCV)*, 740–755.
- [12] Dosovitskiy, A., Beyer, L., Kolesnikov, A., et al. (2021). An image is worth 16x16 words: Transformers for image recognition at scale. *International Conference on Learning Representations (ICLR)*.
- [13] Rahman, A., & Miah, M. (2017). Automatic detection and classification of plant leaf diseases. *International Journal of Information Technology and Computer Science (IJITCS)*, 9(11), 1–8.
- [14] Zhang, Y., Wu, L., & Zhao, G. (2019). Plant leaf disease recognition and classification using improved convolutional neural network. *Frontiers in Plant Science*, 10, 1582.
- [15] Singh, S., Mohapatra, S. K., & Patra, P. K. (2015). Application of neural networks in disease classification and crop management. *Information Processing in Agriculture*, 2(1), 1–8.